

# Deep Learning

## Episode 0

# ML recap. Adaptive optimization

Maxim Borisyak, Alexander Panin, Andrey Ustyuzhanin



Yandex  
Data Factory

LAMBDA



British Hedgehog  
Preservation Society



# Linear Regression

Model:

$$X \longrightarrow Wx + b \longrightarrow Y^{\text{pred}}$$

Objective function:

$$L = \sum_i (y_i - y_i^{\text{pred}})^2$$

Optimization (exact):

$$w = (X^T X)^{-1} X^T y$$

# Linear Regression

Model:

$$X \longrightarrow Wx + b \longrightarrow Y^{\text{pred}}$$

Objective function:

$$L = \sum_i (y_i - y_i^{\text{pred}})^2$$

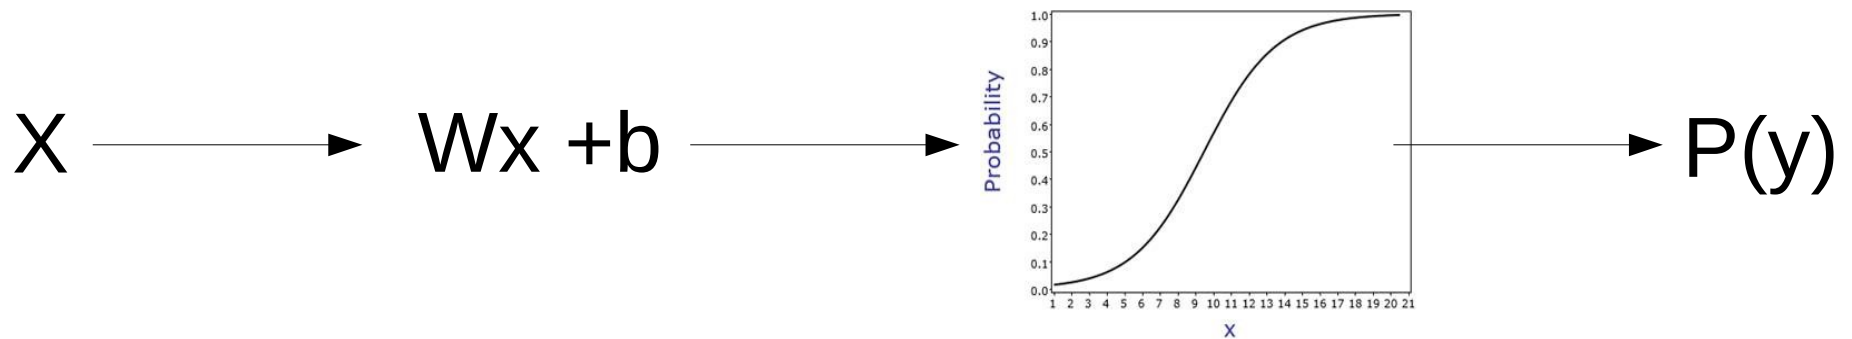
Optimization (iterative):

$$w_0 \leftarrow 0$$

$$w_{i+1} \leftarrow w_i - \alpha \frac{\delta L}{\delta W}$$

$$\frac{\delta L}{\delta W} = \sum_i -2x(y_i - (wx_i + b))$$

# Logistic Regression

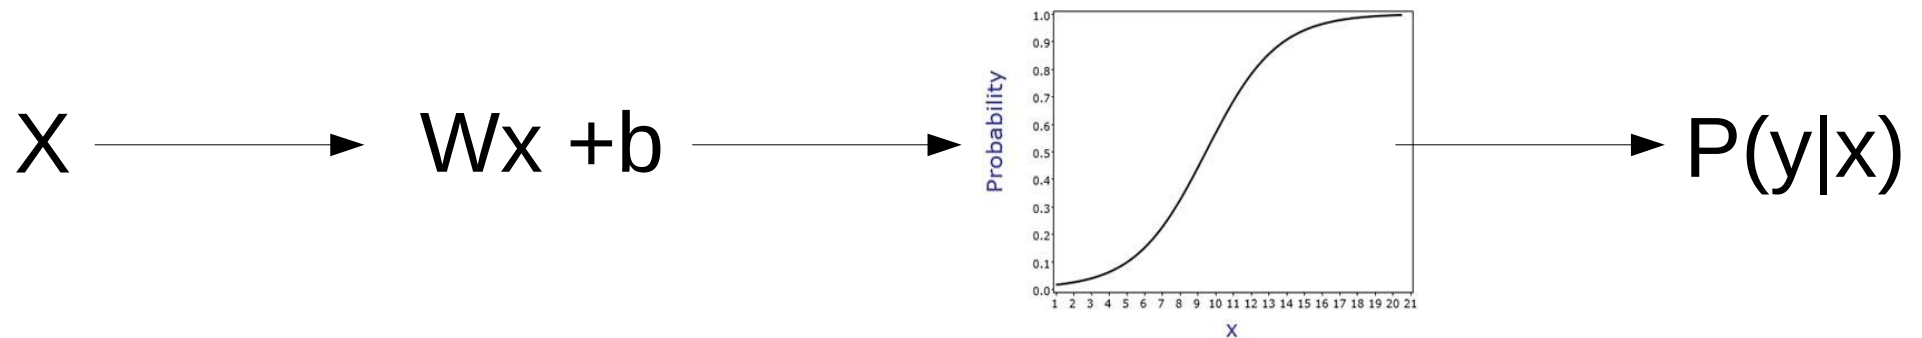


$$P(y) = \sigma(Wx + b)$$

Objective function ?

# Logistic Regression

Model:



Objective function:

$$L = - \sum_i y \log P(y|x) + (1 - y) \log (1 - P(y|x))$$

Optimization (iterative):

You guessed it!

# Logistic Regression

Model:

$$\begin{array}{c} X \longrightarrow \begin{array}{l} a_{[y=a]} = W_a x + b_a \\ a_{[y=b]} = W_b x + b_b \\ a_{[y=c]} = W_c x + b_c \end{array} \longrightarrow \frac{e^{a_{[y=class]}}}{\sum_j e^{a_{[y=j]}}} \longrightarrow \begin{array}{l} P(y=a|X) \\ P(y=b|X) \\ P(y=c|X) \end{array} \end{array}$$

Objective function:


$$L = - \sum_i \sum_{class} [y_i = class] \log P(y_i = class | x)$$

# Gradient descent

Update:

$$w_{i+1} \leftarrow w_i - \alpha \frac{\delta L}{\delta W}$$

- $\alpha$  – learning rate  $\alpha \ll 1$
- $L$  – loss function



Can we do better?


# Newton-Raphson

Parameter update

$$w_{i+1} \leftarrow w_i - \alpha H_L^{-1} \frac{\delta L}{\delta W}$$

Hessian:

$$\mathbf{H} = \begin{bmatrix} \frac{\partial^2 f}{\partial x_1^2} & \frac{\partial^2 f}{\partial x_1 \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_1 \partial x_n} \\ \frac{\partial^2 f}{\partial x_2 \partial x_1} & \frac{\partial^2 f}{\partial x_2^2} & \cdots & \frac{\partial^2 f}{\partial x_2 \partial x_n} \\ \vdots & \vdots & \ddots & \vdots \\ \frac{\partial^2 f}{\partial x_n \partial x_1} & \frac{\partial^2 f}{\partial x_n \partial x_2} & \cdots & \frac{\partial^2 f}{\partial x_n^2} \end{bmatrix}.$$



Red: Newton-Raphson  
Green: gradient descent

Any drawbacks?



# Stochastic gradient descent


Loss function is mean over all data samples.

Approximate with 1 or few random samples.

Update:

$$w_{i+1} \leftarrow w_i - \alpha \frac{E \delta L}{\delta W}$$

- E – expectation
- Learning rate should decrease



# SGD with momentum

Idea: move towards “overall gradient direction”,  
Not just current gradient.

$$w_0 \leftarrow 0 ; v_0 \leftarrow 0$$

$$v_{i+1} \leftarrow \alpha \frac{\delta L}{\delta W} + \mu v_i$$

$$w_{i+1} \leftarrow w_i - v_{i+1}$$

Helps for noisy gradient / canyon problem



# SGD with momentum

Idea: move towards “overall gradient direction”,  
Not just current gradient.

$$w_0 \leftarrow 0; v_0 \leftarrow 0$$

$$v_{i+1} \leftarrow \alpha \frac{\delta L}{\delta W} + \mu v_i$$

$$w_{i+1} \leftarrow w_i - v_{i+1}$$



# AdaGrad

Idea: decrease learning rate individually for each parameter in proportion to sum of it's gradients so far.

$$G_t = \sum_{\tau=1}^t \left[ \frac{\delta L}{\delta w_{\tau}} \right]^2$$

“Total update path length”  
(for each parameter)

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{G_t + \epsilon}} \frac{\delta L}{\delta w_i}$$


# RMSProp

Idea: make sure all gradient steps have approximately same magnitude (by keeping moving average of magnitude)

$$ms_{t+1} = \gamma \cdot ms_t + (1 - \gamma) \left\| \frac{\delta L}{\delta w_{t+1}} \right\|^2$$

$$w_{t+1} = w_t - \frac{\eta}{\sqrt{ms + \epsilon}} \frac{\delta L}{\delta w_i}$$

# Alltogether



# Moar stuff

## **Without Hessian**

- Adadelata ~ adagrad with window
- Adam ~ rmsprop + momentum
  - Nesterov-momentum
  - Hessian-free (narrow)
  - Conjugate gradients

## **Estimate inverse Hessian**

- BFGS
- L-BFGS
- \*\*\*\*-BFGS

# Regularization (weight)

General idea:

$$L_{new} = L + reg$$

performance = how\_i\_fit\_data + how\_reasonable\_i\_am

L2 regularizer

$$L_{new} = L + \|\theta\|_2^2 = L + \sum_i \theta_i^2$$

linear models:  $\theta = \{w, b\}$

- a.k.a. weight decay
- a.k.a. Tikhonov regularizer
- a.k.a. normal prior on params



# Regularization (weight)

L2 regularizer

$$L_{new} = L + \sum_i \theta_i^2$$

L1 regularizer

$$L_{new} = L + \sum_i |\theta_i|$$

Difference between L1, L2?


Any other way to regularize?

# Regularization(other)


- Distort input
- Distort weights
- Additional objective
- Domain-specific stuff
- Moar data :)
- etc.

Most are domain- or model-specific

# Nonlinear dependencies



What we have



What we want

- How to get that?

# Nuff

**Let's go implement that!**

